



# Book 1: Product (MVP Spec)



Cynqra Product Book, v0.1. This book answers: what are we building first, precisely, and how do we know it works? Owner: CPO. This book changes constantly. Every change is versioned at the bottom. Concepts used here (Worker, Verification, Transition and so on) are defined in Book 0.

1. Product thesis

2. Target user

3. Scope

P0: must ship in the MVP

P1: should ship if P0 holds

Explicit non goals for the MVP

4. User journeys

5. Progressive autonomy

6. Metrics

7. Top product risks

8. Open questions

9. The entry point: where objectives are born

## 1. Product thesis

The MVP proves the first transition of the core loop (Book 0, §4 and §26):

**A founder describes a software product. Cynqra assembles a fixed template AI organization and delivers a deployed, working web application, with the**

**founder approving everything above low risk.**

The demo in one sentence: *"I typed an objective, approved an org and a plan, answered a handful of decisions, and two weeks later I had a deployed product I could export and walk away with."*

## 2. Target user

| Persona                                     | Description                                                                                  | Priority       |
|---------------------------------------------|----------------------------------------------------------------------------------------------|----------------|
| <b>A. Solo technical founder</b>            | Can read code, has shipped before, wants leverage. Tolerates rough UX if the output is real. | <b>Primary</b> |
| <b>B. Non technical founder with budget</b> | Wants a product built without hiring a team. Needs heavy guardrails.                         | Phase 2        |

Why A first: code output is verifiable (tests and builds), technical users forgive rough edges, and their feedback is precise. Persona B arrives when verification and UX mature.

## 3. Scope

### P0: must ship in the MVP

| # | Requirement                                         | Done when                                                                                                                                       |
|---|-----------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------|
| 1 | Founder account and company creation                | A founder creates an account and a company in under 5 minutes                                                                                   |
| 2 | Objective intake (structured wizard, not free chat) | A free text objective converts to a structured objective (product, target, outcomes, success criteria, constraints) and the founder confirms it |
| 3 | Fixed org template                                  | The system instantiates one template: CTO worker, PM worker, 2 engineer workers,                                                                |

| # | Requirement                         | Done when                                                                                                                                                              |
|---|-------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|   |                                     | QA and verification function. No dynamic org synthesis in the MVP                                                                                                      |
| 4 | Organizational graph                | Every task, worker and decision is queryable by owner, dependency and approver, and rendered as a live tree in the UI                                                  |
| 5 | Task generation and dependencies    | Objective → workstreams → tasks with the 12 task fields from Book 0 §13, dependency graph included                                                                     |
| 6 | Worker runtime with sandboxed tools | Workers execute in isolated containers with allowlisted tools only: repo, file editor, shell, test runner, deploy target                                               |
| 7 | Coordination protocols (MVP subset) | Blocker, Handoff, Escalation and Approval work end to end, and a cross worker dependency resolves without founder involvement at least once per run                    |
| 8 | Verification pipeline               | LOW risk → automated checks (build and tests). MEDIUM → independent second model review. HIGH → founder approval. The UI visibly separates "completed" from "verified" |
| 9 | Approval inbox                      | The founder sees a decision queue with problem, evidence, recommendation, cost, risk and confidence, plus Approve / Reject / Request more evidence                     |

| #  | Requirement                         | Done when                                                                                            |
|----|-------------------------------------|------------------------------------------------------------------------------------------------------|
| 10 | Budget tracking and circuit breaker | There is a per company budget cap with a hard stop and a founder alert when the cap is hit           |
| 11 | Event log and audit trail           | Every action is queryable: who acted, authority, context, tools, outcome (Book 0 §20)                |
| 12 | Delivery and export                 | Working deployed URL plus full repo and knowledge export. The portability promise holds from day one |

## P1: should ship if P0 holds

Worker performance stats page. Evolution dashboard (view only recommendations). Experiment engine lite. Dashboard polish. Model fallback chain.

## Explicit non goals for the MVP

Dynamic org synthesis. Human workers in the loop. Intelligence marketplace. Non software company types. Autonomous high risk actions. Broad external integrations. Mobile app.

## 4. User journeys

1. **Onboard to objective.** Sign up → create company → objective wizard → review the structured objective → set budget cap and risk tolerance → confirm.
2. **Org proposal to approval.** Cynqra presents the fixed org, plan, milestones and projected cost → founder approves → execution begins.
3. **Daily loop.** The founder opens the Decision dashboard (target: under 15 minutes a day) → clears the approval queue → reviews flagged evidence → can inspect any worker's work.

4. **Delivery.** Verification report per milestone → founder accepts → deployed URL plus export bundle → transition record written.

## 5. Progressive autonomy

| Level | Name                                   | Behavior                                                   |
|-------|----------------------------------------|------------------------------------------------------------|
| L0    | Observation                            | Cynqra recommends everything, executes nothing             |
| L1    | Recommendation plus low risk execution | Auto executes LOW risk actions only                        |
| L2    | Bounded autonomy                       | MEDIUM risk actions auto execute within budget and limits  |
| L3    | Expanded autonomy                      | Authority expands per worker, earned by performance record |

New companies start at L1. Level changes are founder decisions and are always logged.

## 6. Metrics

**North star:** founder interventions per week (Book 0 §16), trending down without outcome quality dropping.

**Input metrics:** task success rate (verified first pass %), verification pass rate, escalation rate per 100 tasks, cost per verified milestone, time from objective to working product, rollback rate, export success rate.

**Kill criteria for the MVP:** if after 3 full company runs, under 40% of tasks pass verification without rework, or the cost per run exceeds what a freelance team would charge, we stop and reassess the architecture before scaling.

## 7. Top product risks

| Risk                                                      | Mitigation                                                                       |
|-----------------------------------------------------------|----------------------------------------------------------------------------------|
| Verification fails on non code work (research, specs, UX) | The MVP scope is limited to software. Non code outputs are always MEDIUM risk or |

| Risk                                                            | Mitigation                                                                                                      |
|-----------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------|
|                                                                 | higher, so they get independent review or the founder                                                           |
| Founder drowns in escalations, which contradicts the north star | Escalation budget per day, batching, severity filters, measured weekly                                          |
| Coordination token cost explodes                                | Protocol messages are structured, not free chat. Context layering per Book 0 §19. Cost circuit breaker (P0 #10) |
| Trust: a founder will not grant authority to an unproven system | Progressive autonomy ladder (§5), full audit trail, exportability                                               |

## 8. Open questions

1. What is the narrowest product category for the first runs? Landing pages, CRUD apps, internal tools? Our recommendation: CRUD style web apps, because they give the richest automated verification.
2. Flat per run pricing or subscription plus usage? Book 3 models cost first.
3. How much founder input during planning: one shot approval or interactive?

## 9. The entry point: where objectives are born

Raised in the bench review: the objective gets typed somewhere, and whoever owns that moment owns the customer relationship.

1. **We own the moment of intent.** Cynqra is not a chatbot that receives objectives. It is the place where an idea becomes a company. The objective wizard is the front door: structured, guided, and backed by the Company Genome (what must be true for this to succeed). A generic chat box cannot replicate it, because a chat box produces a conversation. Cynqra produces a company.
2. **The first session is the product.** A founder types one sentence and, in the same sitting, sees a structured objective, a proposed organization, a plan, a projected cost, and the first three decisions worth their judgment.

3. **Distribution.** We meet founders where ideation happens (communities, accelerators, founder content) and we make the first run irresistible. The door is the demo.

**The scale invariant (design requirement).** The founder surface stays at roughly **three decisions a day at any org size**. Seven workers or seven hundred, the inbox compresses: hierarchy, escalation budgets and authority policies decide what reaches the founder. The org graph is not only for organizing workers. It is for protecting the founder's attention. The dashboard shows state, constraints and three decisions. Never a feed.

**The habit.** A door wins the first session. A rhythm keeps the company. The founder's daily loop is designed as a morning ritual: three decisions, the constraint of the week, and one outcome report. Weekly, the Evolution review proposes organizational changes. The founder does not manage the company. They keep a standing appointment with it.

*Version history: v0.3. §9 extended with the habit loop (bench round 3). v0.2: added §9, the entry point and the scale invariant. v0.1: first draft, derived from Book 0 §26 and §27.*

# Book 1: Product — Production Build Addendum v0.4

Book 1 remains the Product/MVP source of truth.

Purpose: close the one confirmed Book 0/Product contradiction and convert existing P0 requirements into testable acceptance contracts. No new product capability is introduced.

## Binding correction: Progressive Autonomy

This addendum supersedes Book 1 §5 only where the existing wording conflicts with Book 0 §12. The Doctrine remains authoritative.

| Level | Name                       | Binding behavior                                                                                                                                                                                        |
|-------|----------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| L0    | Observation                | Cynqra recommends; it executes no actions.                                                                                                                                                              |
| L1    | Bounded low-risk execution | Worker may execute only LOW-risk actions explicitly permitted by policy, within budget, tool, environment, and rate limits.                                                                             |
| L2    | Delegated bounded autonomy | Worker may execute only defined MEDIUM-risk action classes for which authority has been explicitly delegated and the applicable review condition is satisfied. L2 is not blanket MEDIUM-risk authority. |
| L3    | Expanded autonomy          | Authority may expand only by explicit action-class policy and demonstrated reliability. L3 never overrides HIGH-risk human approval or MVP prohibitions.                                                |

Permanent MVP safety boundary: HIGH-risk actions remain human gated at every autonomy level. No autonomy level permits autonomous legal or financial commitments, money movement, external messages, objective changes, budget changes, authority-boundary changes, or other actions prohibited by Book 0/Book 3.

## MVP acceptance contract

The existing P0 scope remains unchanged. The following are implementation acceptance conditions for the existing requirements, not new product features.

| Requirement        | Production acceptance condition                                                                                                                                             |
|--------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Objective intake   | Structured objective is persisted with product, target, outcomes, success criteria, constraints, priority and version; founder confirmation is recorded.                    |
| Fixed organization | Exactly the MVP template is instantiated. Dynamic organization synthesis is not invoked.                                                                                    |
| Task execution     | Every task has owner, inputs, expected output, dependencies, tools, budget, deadline, verification method, risk tier and status.                                            |
| Coordination       | Blocker, Handoff, Escalation and Approval are structured records/events; at least one cross-worker dependency resolves without founder involvement during a qualifying run. |
| Verification       | Completed and verified are separate states; verdict and evidence are persisted.                                                                                             |
| Budget             | Hard company cap, warning thresholds, breaker state and founder notification are observable and auditable.                                                                  |
| Audit              | Important actions reconstruct actor, authority, context, intelligence, tools, outcome and resulting event.                                                                  |
| Delivery           | Deployment record, health/smoke verification, live URL, and export manifest are present.                                                                                    |

## Metrics and kill criteria clarification

The existing North Star and kill criteria remain unchanged. For engineering measurement, first-pass verification must be reported together with defect escape, false acceptance, false rejection, verification latency, verification cost and rework rate. These supporting measures do not replace the existing kill criteria.

Founder intervention must never be optimized in isolation. A lower intervention count is an improvement only when verified outcome quality remains within the accepted range and there is no evidence of suppressed escalation.

## **Open Product decisions**

The existing Book 1 open questions remain open unless Book 4 changes their status. The first validation runs may use CRUD-style web apps as the current recommendation, but this addendum does not convert that recommendation into a company-wide permanent decision.

Pricing remains deferred until the cost model has real data. This does not block engineering of the MVP execution path.

Planning interaction mode remains a product decision. The implementation must support the existing founder approval step and preserve the ability to evolve the interaction without changing the underlying authority model.